

1. Data Loading and Preprocessing:

- Load the dataset from a CSV file.
- Clean the dataset by removing rows with missing values.
- Tokenize the input texts using the BERT tokenizer.
- Convert the labels to tensors.

2. Model Definition:

- Define a custom neural network model called BertCNNClassifier.
- The model architecture combines BERT embeddings with a convolutional neural network (CNN).
- BERT embeddings are obtained from the pre-trained BERT model.
- CNN layers are applied to capture features from the BERT embeddings.
- The final layer is a fully connected layer for classification.

3. Model Training:

- Split the data into training and validation sets.
- Define the loss function (CrossEntropyLoss) and optimizer (Adam) with weight decay.
- Train the model for a specified number of epochs.
 - In each epoch:
 - Iterate over the training data in batches.
 - Forward pass: Compute the output predictions using the model.
 - Compute the loss between the predictions and the actual labels.
 - Backpropagate the gradients and update the model parameters.
 - Monitor and record the training loss.
 - Evaluate the model on the validation set after each epoch:
 - Iterate over the validation data in batches.
 - Forward pass: Compute the output predictions using the model.

- Compute the validation loss.
- Monitor and record the validation loss.

4. Model Evaluation and Saving:

- Save the model state with the best validation accuracy.
- Optionally, you can load the saved model state to make predictions on new data.

Summary:

This code implements a text classification model that combines BERT embeddings with a CNN architecture. It trains the model on a labeled dataset, monitors the training progress, and evaluates its performance on a validation set. Finally, it saves the best-performing model for future use.